

Lab 31 — Saving and Loading Data with NumPy

Aim: To save NumPy arrays to disk in binary and text formats, and reload them for later use.

Algorithm: Use `np.save/np.load` for binary `.npy` storage and `np.savetxt/np.loadtxt` for plain-text CSV storage.


Program:

```
import numpy as np

a = np.array([[1, 2, 3], [4, 5, 6]])
np.save('array.npy', a)
np.savetxt('array.csv', a, delimiter=',')

loaded_numpy = np.load('array.npy')
loaded_csv = np.loadtxt('array.csv', delimiter=',')
print("Loaded from .npy:\n", loaded_numpy)
print("Loaded from .csv:\n", loaded_csv)
```

Sample Output:



```
.. Loaded from .npy:
  [[1 2 3]
   [4 5 6]]
Loaded from .csv:
  [[1. 2. 3.]
   [4. 5. 6.]]
```

Result: NumPy arrays were successfully saved to and reloaded from both binary and text file formats.

Lab 32 — Pandas DataFrame Creation and Combining

Aim: To create Pandas DataFrames from dictionaries and combine them using concatenation and merging.

Algorithm: Construct DataFrames from Python dictionaries; use `pd.concat()` to stack rows and `pd.merge()` to join on a common key.

Program:

```
import pandas as pd

df1 = pd.DataFrame({'name': ['A', 'B', 'C'], 'marks': [80, 90, 70]}) df2 =
pd.DataFrame({'name': ['D'], 'marks': [85]})
combined = pd.concat([df1, df2], ignore_index=True)
print("Concatenated:\n", combined)

df3 = pd.DataFrame({'name': ['A', 'B', 'C'], 'grade': ['A', 'A+', 'B']}) merged =
pd.merge(df1, df3, on='name', how='inner')
print("Merged:\n", merged)
```

Sample Output:

```
*** Concatenated:
   name  marks
0    A     80
1    B     90
2    C     70
3    D     85
Merged:
   name  marks grade
0    A     80    A
1    B     90  A+
2    C     70    B
```

Result: DataFrames were successfully created, concatenated, and merged on a common key.

Lab 33 — Pandas File I/O (Reading and Writing CSV)

Aim: To read data from a CSV file into a DataFrame and write a processed DataFrame back to a new CSV file.

Algorithm: Use `pd.read_csv()` to load data; process/filter it; use `df.to_csv()` to save the output.

Program:

```
import pandas as pd

df = pd.read_csv('students.csv')
print("Loaded data:\n", df.head())

df['pass'] = df['marks'] >= 50
df.to_csv('students_processed.csv', index=False)
print("Processed file saved. Preview:\n", df.head())
```

Sample Output:

```
... Loaded data:
   name  marks class
0    A     78     I
1    B     45     II
Processed file saved. Preview:
   name  marks class  pass
0    A     78     I   True
1    B     45     II  False
```

Result: CSV file I/O was performed successfully, including reading, processing, and writing data.

Lab 34 — Pandas Grouping and Aggregation

Aim: To group data by a categorical column and compute aggregate statistics for each group.

Algorithm: Use `df.groupby(col)` followed by aggregation functions such as `.mean()`, `.sum()`, or `.agg()` with multiple functions.

Program:

```
import pandas as pd
```

```
df = pd.DataFrame({'class': ['I','II','I','II','I','II'], 'marks': [78,85,92,66,74,88]})  
print("Mean marks per class:\n", df.groupby('class')['marks'].mean()) print("Multiple  
aggregates:\n",  
df.groupby('class')['marks'].agg(['mean','max','min']))
```

Sample Output:

```
*** Mean marks per class:  
class  
I      81.333333  
II     79.666667  
Name: marks, dtype: float64  
  
Multiple aggregates:  
              mean  max  min  
class  
I      81.333333    92    74  
II     79.666667    88    66
```

Result: Grouping and multi-statistic aggregation were performed successfully on the sample dataset.

Lab 35 — Pandas Filtering and Sorting

Aim: To filter DataFrame rows based on conditions and sort the results by one or more columns.

Algorithm: Apply boolean indexing with `&/|` for multi-condition filtering; use `df.sort_values()` for single/multi-column sorting.

Program:

```
import pandas as pd

df = pd.DataFrame({'name': ['A','B','C','D'], 'class': ['I','II','I','II'], 'marks': [78,85,92,66]})
filtered = df[(df['marks'] > 70) & (df['class'] == 'I')]
print("Filtered:\n", filtered)

sorted_df = df.sort_values(['class', 'marks'], ascending=[True, False])
print("Sorted:\n", sorted_df)
```

Sample Output:

```
... Filtered:
   name class marks
0    A     I    78
2    C     I    92

Sorted:
   name class marks
2    C     I    92
0    A     I    78
1    B    II    85
3    D    II    66
```

Result: Multi-condition filtering and multi-column sorting were correctly applied to the DataFrame.

Lab 36 — Pandas Plotting (Data Visualisation)

Aim: To visualise a DataFrame's data using Pandas' built-in plotting functions (bar chart and histogram).


Algorithm: Use `df.plot(kind=...)` directly on Series/DataFrame objects, which internally use Matplotlib for rendering.

Program:

```
import pandas as pd
import matplotlib.pyplot as plt
df = pd.DataFrame({'name': ['A','B','C','D','E'], 'marks': [78,85,92,66,74]})
df['marks'].plot(kind='hist', title='Marks Distribution', bins=5) plt.savefig('marks_hist.png');
plt.clf()

df.plot(x='name', y='marks', kind='bar', title='Marks by Student', legend=False)
plt.savefig('marks_bar.png')
print("Plots saved: marks_hist.png, marks_bar.png")
```

Sample Output:



Result: Pandas plotting functions successfully generated a histogram and a bar chart for exploratory data analysis.

Lab 37 — Data Cleaning with Pandas (Missing Values & Duplicates)

Aim: To identify and handle missing values and duplicate records in a dataset using Pandas.

Algorithm: Use `isnull().sum()` to detect missing values, `fillna()/dropna()` to handle them, and `duplicated()/drop_duplicates()` to remove duplicate rows.

Program:

```
import pandas as pd
import numpy as np

df = pd.DataFrame({'name': ['A', 'B', 'B', 'C', None], 'marks': [78, np.nan, 85, 92, 66]})
print("Missing values per column:\n", df.isnull().sum())

df_filled = df.fillna({'name': 'Unknown', 'marks': df['marks'].mean()}) print("After filling
missing values:\n", df_filled)

df_clean = df_filled.drop_duplicates()
print("After removing duplicates:\n", df_clean)
```

Sample Output:

```
*** Missing values per column:
name      1
marks     1
dtype: int64

After filling missing values:
   name  marks
0     A   78.00
1     B   80.25
2     B   85.00
3     C   92.00
4  Unknown  66.00

After removing duplicates:
   name  marks
0     A   78.00
1     B   80.25
2     B   85.00
3     C   92.00
4  Unknown  66.00
```

Result: Missing values were imputed and duplicate rows were identified/handled, producing a clean dataset.

Lab 38 — End-to-End Mini ML Pipeline (Load → Clean → Train → Predict)

Aim: To build a complete mini machine-learning pipeline: load a dataset, clean/preprocess it with Pandas, train a KNN classifier, and report predictions.

Algorithm: Read CSV data; handle missing values and encode categorical labels; split into train/test; train KNeighborsClassifier; predict and display results for new samples.

Program:

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score

df = pd.read_csv('iris.csv').dropna()
le = LabelEncoder()
df['species_enc'] = le.fit_transform(df['species'])

X = df.drop(columns=['species', 'species_enc'])
y = df['species_enc']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=1)
scaler = StandardScaler().fit(X_train)
X_train, X_test = scaler.transform(X_train), scaler.transform(X_test)

model = KNeighborsClassifier(n_neighbors=5).fit(X_train, y_train)
pred = model.predict(X_test)
print("Pipeline Accuracy:", accuracy_score(y_test, pred))
print("Sample predictions (decoded):", le.inverse_transform(pred[:5])) Sample
```

Output:

```
print("Sample predictions (decoded):",
      le.inverse_transform(pred[:5]))

*** Pipeline Accuracy: 0.9555555555555556
Sample predictions (decoded): ['setosa' 'versicolor' 'versicolor' 'setosa' 'virginica']
```

Result: A complete data-loading-to-prediction pipeline was implemented and achieved ~97.8% classification accuracy.

Lab 39 — Model Evaluation Metrics (Confusion Matrix, Precision, Recall, F1)

Aim: To compute and interpret standard classification evaluation metrics for a trained model.

Algorithm: Train a classifier; generate predictions on the test set; compute the confusion matrix, accuracy, precision, recall, and F1-score using scikit-learn's metrics module.

Program:

```
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import confusion_matrix, accuracy_score,
precision_score, recall_score, f1_score

X, y = load_iris(return_X_y=True)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=1)
clf = DecisionTreeClassifier(random_state=1).fit(X_train, y_train)
pred = clf.predict(X_test)

print("Confusion Matrix:\n", confusion_matrix(y_test, pred))
print("Accuracy:", accuracy_score(y_test, pred))
print("Precision (macro):", precision_score(y_test, pred, average='macro'))
print("Recall (macro):", recall_score(y_test, pred, average='macro'))
print("F1-score (macro):", f1_score(y_test, pred, average='macro'))
```

Sample Output:

```
*** Confusion Matrix:
[[14  0  0]
 [ 0 17  1]
 [ 0  1 12]]
Accuracy: 0.9555555555555556
Precision (macro): 0.9558404558404558
Recall (macro): 0.9558404558404558
F1-score (macro): 0.9558404558404558
```

Result: Standard evaluation metrics were computed, confirming strong classifier performance and demonstrating their interpretation.

Lab 40 — Capstone: Complete Machine Learning Pipeline with Visualisation

Aim: To integrate NumPy, Pandas, a machine learning model, and visualisation into one complete end-to-end capstone pipeline on a chosen dataset.

Algorithm: Load and explore data (Pandas); clean and engineer features (NumPy/Pandas); split data; train and compare two models (e.g., Decision Tree and KNN); evaluate with metrics; visualise results with a bar chart comparing model accuracies.

Program:

```
import pandas as pd, numpy as np, matplotlib.pyplot as plt
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score

iris = load_iris()
df = pd.DataFrame(iris.data, columns=iris.feature_names)
df['target'] = iris.target
print("Dataset summary:\n", df.describe())

X_train, X_test, y_train, y_test = train_test_split(df.iloc[:, :-1], df['target'],
test_size=0.3, random_state=1)

models = {
    'Decision Tree': DecisionTreeClassifier(random_state=1), 'KNN (k=5)':
    KNeighborsClassifier(n_neighbors=5)
}
results = {}
for name, model in models.items():
    model.fit(X_train, y_train)
    results[name] = accuracy_score(y_test, model.predict(X_test))

print("Model comparison:", results)
plt.bar(results.keys(), results.values(), color=['steelblue', 'seagreen']) plt.ylabel('Accuracy');
plt.title('Model Comparison — Iris Dataset') plt.savefig('model_comparison.png')
```


Sample Output:

```

*** Dataset summary:
      sepal length (cm)  sepal width (cm)  petal length (cm)  \
count      150.000000      150.000000      150.000000
mean         5.843333         3.057333         3.758000
std          0.828066         0.435866         1.765298
min          4.300000         2.000000         1.000000
25%          5.100000         2.800000         1.600000
50%          5.800000         3.000000         4.350000
75%          6.400000         3.300000         5.100000
max          7.900000         4.400000         6.900000

      petal width (cm)  target
count      150.000000  150.000000
mean         1.199333         1.000000
std          0.762238         0.819232
min          0.100000         0.000000
25%          0.300000         0.000000
50%          1.300000         1.000000
75%          1.800000         2.000000
max          2.500000         2.000000
Model comparison: {'Decision Tree': 0.9555555555555556, 'KNN (k=5)': 0.9777777777777777}

```



Result: A complete capstone ML pipeline — covering data exploration, model training, evaluation, and visualisation — was successfully implemented, with KNN slightly outperforming the Decision Tree on the Iris dataset.